

TECH CHALLENGE

I created a Repo and cloned this locally

[PYTHON SCRIPT](#)

[SERVER](#)

[DOCKER](#)

[K8s](#)

[MINIKUBE](#)

[HORIZONTAL POD AUTOSCALER](#)

PYTHON SCRIPT

In this repo I firstly created the python script and named this `print_timestamp.py`

```
Python
import time

# ts stores the time in seconds
ts = time.time()

# print the current timestamp
print(ts)
```

Running with `python3 print_timestamp.py` allowed me to test that the script was successful and ensured that the task of attaining console type output can be cross-checked and referenced to the output of the kubernetes cluster once deployed.

The naming format utilising an underscore was retained as important so that all understood the name to be referenced.

```
Unset
> python3 print_timestamp.py
1716319109.759373
```

SERVER

To ensure a small image size, it was decided to use the built-in Python server to return the time.

Python

```
#import need libraries
import time
from http.server import BaseHTTPRequestHandler, HTTPServer

#Class to be used
class RequestHandler(BaseHTTPRequestHandler):
    def do_GET(self):
        # Get the current timestamp
        ts = time.time()

        # Send response status code and headers
        self.send_response(200)
        self.send_header('Content-type', 'text/plain')
        self.end_headers()

        # Send the timestamp as a response
        self.wfile.write(str(ts).encode())

#Run the server and allow response
def run(server_class=HTTPServer, handler_class=RequestHandler, port=4430):
    server_address = ('', port)
    httpd = server_class(server_address, handler_class)
    print(f'Starting httpd server on port {port}...')
    httpd.serve_forever()

if __name__ == '__main__':
    run()
```

This was run in one terminal and in a second running `curl 127.0.0.1:4430` calls to the server and returned `1716319485.656218%`

This was different to the original output, adding %. After research, this was due to curl adding this to indicate the end of line as python was not returning a new line.

New line was added to the script and not the % was no longer seen

```
self.wfile.write((str(ts) + '\n').encode())
```

DOCKER

Now we need to create the docker image to be used in K8s.

```
Unset
# Use a lightweight Python image
FROM python:3.9-slim

# Set the working directory in the container
WORKDIR /app

# Copy the Python script into the container
COPY print_timestamp.py .

# Command to run the Python script
```

```
CMD ["python", "print_timestamp.py"]
```

Running `docker build -t print_timestamp:latest .` this will build the image locally for us to run.

Unset

```
> docker images
REPOSITORY          TAG          IMAGE ID        CREATED         SIZE
print_timestamp     latest      6b323af908dc   55 minutes ago 152MB
```

Running `docker run -p 4430:4430 print_timestamp` this will now run the image locally to allow testing

Unset

```
> docker ps -a
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS
780f9746a412   print_timestamp "python print_timest..." 35 seconds ago Up 34 seconds 0.0.0.0:4430->4430/tcp
```

Again running the curl command will call the running server that is open on port 4430, this giving the expected and same reply as the original script output.

K8s

Now to look at the K8s section and to allow local testing, I used minikube. This can be built from [here](#)

Once installed we can start minikube with `minikube start`

Running `kb get namespace` will show minikube running.

Unset

```
> kb get namespaces
NAME           STATUS    AGE
default        Active   100m
kube-node-lease Active   100m
kube-public    Active   100m
kube-system    Active   100m
```

We created a deployment yaml.

Unset

```
apiVersion: apps/v1
kind: Deployment
```

```

metadata:
  name: print_timestamp
spec:
  replicas: 1
  selector:
    matchLabels:
      app: print_timestamp
  template:
    metadata:
      labels:
        app: print_timestamp
    spec:
      containers:
        - name: print_timestamp
          image: print-timestamp:latest
          ports:
            - containerPort: 4430

```

On trying to run this (`kb apply -f k8s/deployment.yaml`) the underscore I had retained in the naming convention bit me with the error.

Unset

```

The Deployment "print_timestamp" is invalid: spec.template.spec.containers[0].name: Invalid value:
"print_timestamp": a lowercase RFC 1123 label must consist of lower case alphanumeric characters or
'-', and must start and end with an alphanumeric character (e.g. 'my-name', or '123-abc', regex used
for validation is '[a-z0-9]([-a-z0-9]*[a-z0-9])?')

```

I changed all naming to change the _ (snake) to - (kebab) to maintain the ease and consistency.

Unset

```

> cp print_timestamp.py print-timestamp.py
> docker build -t print-timestamp:latest .
> docker images
REPOSITORY          TAG          IMAGE ID          CREATED           SIZE
print-timestamp     latest      2cf3cf13630f     57 seconds ago   152MB
print_timestamp     latest      6b323af908dc     2 hours ago      152MB
> docker rmi 6b323af908dc
Untagged: print_timestamp:latest
Deleted: sha256:6b323af908dc2f86dc8c7d70c446683d4dc7b97145b0597a644352b2f6b38bcc

```

I realised this image appeared to be a little chunky and so changed to use alpine as the base.

Unset

```

# Use a lightweight image
FROM alpine:latest

```

```
# Install Python3 and pip
RUN apk add --no-cache python3

# Set the working directory in the container
WORKDIR /app

# Copy the Python script into the container
COPY print-timestamp.py .

# Command to run the Python script
CMD ["python", "print-timestamp.py"]
```

Unset

```
> docker build -t print-timestamp:alpine .
> docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
print-timestamp	alpine	384f39153a36	3 minutes ago	76MB
print-timestamp	latest	2cf3cf13630f	8 minutes ago	152M

This was half the size so much more pleasing, it was tested to ensure it was working correctly.

I now updated the deployment yaml to kebab case

Unset

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: print-timestamp
spec:
  replicas: 1
  selector:
    matchLabels:
      app: print-timestamp
  template:
    metadata:
      labels:
        app: print-timestamp
    spec:
      containers:
        - name: print-timestamp
          image: print-timestamp:latest
          ports:
            - containerPort: 4430
```

Again running

Unset

```
> kb apply -f k8s/deployment.yaml
deployment.apps/print-timestamp created
> kb get pods -n default
```

NAME	READY	STATUS	RESTARTS	AGE
print-timestamp-ff5bb948f-9kpdr	0/1	Running	0	8m21s

We now create the service yaml

Unset

```
apiVersion: v1
kind: Service
metadata:
  name: print-timestamp
spec:
  selector:
    app: print-timestamp
  ports:
    - protocol: TCP
      port: 4430
      targetPort: 4430
  type: ClusterIP
```

We now run and check the services

Unset

```
> kb apply -f k8s/service.yaml
service/print-timestamp created
> kb get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3h5m
print-timestamp	ClusterIP	10.99.34.31	<none>	4430/TCP	5m15s

We now use port forwarding to allow local connection to the cluster

Unset

```
kubectrl port-forward pod/print-timestamp-5b76b58757-84z2k 4430:4430

Forwarding from 127.0.0.1:4430 -> 4430
Forwarding from [::1]:4430 -> 4430
```

Again using curl we are able to gain a time stamp

Unset

```
curl 127.0.0.1:4430
1716331479.8028572
```

With output given in the first terminal that the connection was made and served
Handling connection for 4430

HORIZONTAL POD AUTOSCALER

We create a scaling policy

```
Unset
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: print-timestamp
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: print-timestamp
  minReplicas: 1
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 50
```

We apply this and can now test.
kb apply -f k8s/service.yaml

TBC.....
TESTING and MONITOR
HELM

We are to test the hpa next using hey to ensure autoscaling works.

I deleted minikube to ensure we had a clean env

Ran minikube addons enable metrics-server to ensure we can pull metrics

Ran minikube image load print-timestamp:alpine to ensure image is available

Ran kb port-forward pod/print-timestamp-5b76b58757-8lzsq 4430:4430 to port forward

Tested using curl again

Running hey -z 30s -c 100 <http://127.0.0.1:4430> places load on the cluster and should force scaling